

COMP 760 Week 3: Generative Models Primer I

By Joey Bose and Prakash Panangaden



Generative models

Destroying structure by adding noise to data is straightforward; transforming noise to structured data is *generative modelling*.

Ingredients Of A Generative Model





Image Super Resolution



Text to Speech



Drug Discovery

Ingredients Of A Generative Model





Image Super Resolution




Text to Speech



Drug Discovery

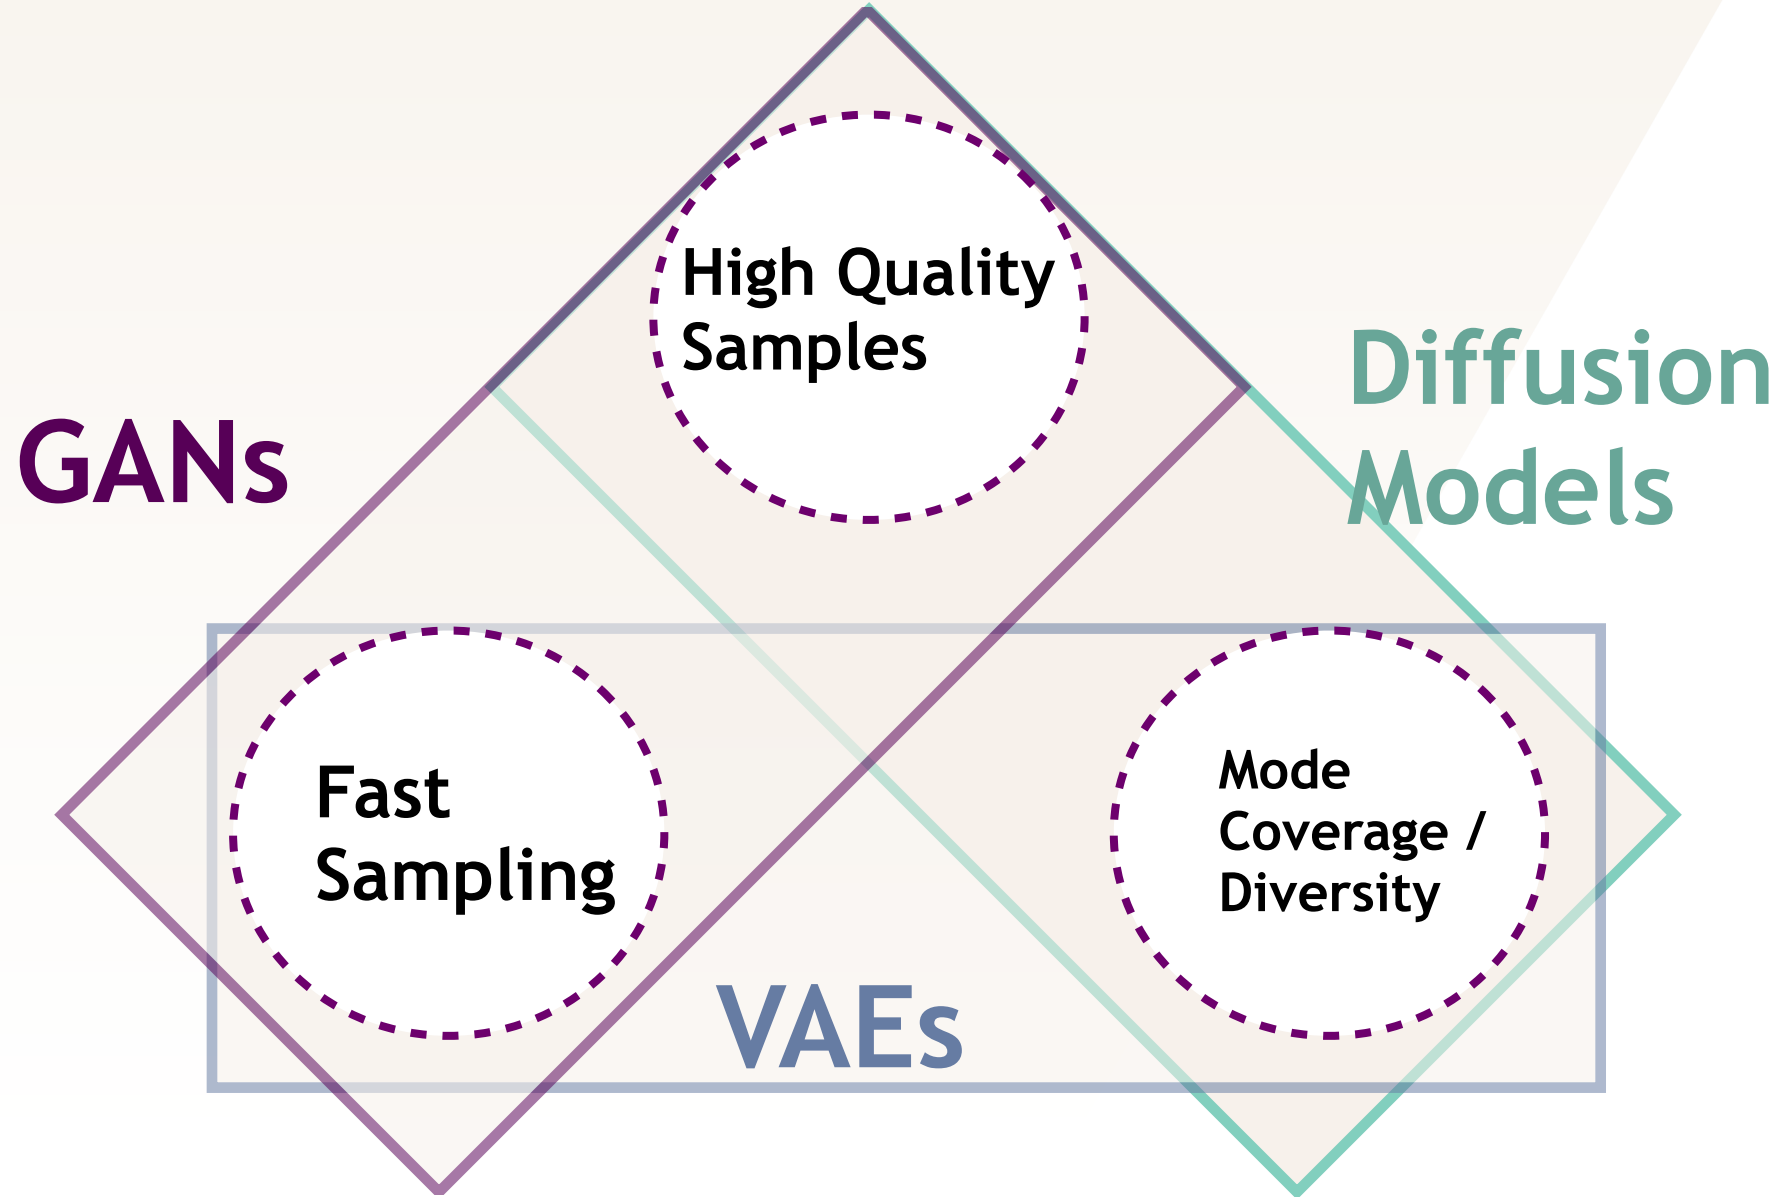
What do we know about the data already?

Ingredients Of A Generative Model




Generative Model Soup

Tradeoffs in a Generative Model





General Setup




Data (Empirical Data Distribution)

General Setup




Data + Noise



We want to model this

$$x \in \mathbb{R}^n \quad \mathcal{D} = \{x_i\} \quad i \in \{1, \dots, N\}$$


Families of Generative Models




Models

● Fully Observed Models

Models that observe all data components directly without introducing any assumptions about unobserved variables.




Families of Generative Models



Models


● Fully Observed Models

Models that observe all data components directly without introducing any assumptions about unobserved variables.




$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Families of Generative Models




Fully Observed Models



$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

- Explicit form for Log-Likelihoods
- Efficient to scale up (e.g. Large Language Models)
- Order Sensitive
- Sequential Generation which can be slow

Families of Generative Models





Models


- Latent Variable Models

Prescribed Models: Assume observed likelihoods plus some observation noise (VAE's, Flows).

Implicit Models: Likelihood Free Models (e.g. GANs).




Families of Generative Models



Latent Variable Models

- Fast Sampling
 - Avoids order dependency; can encode priors
-
- Model Class to Approximate the Posterior may not be rich enough
 - Inversion process from inputs to latent might be hard



Families of Generative Models



Learning Principles

- Maximum Likelihood
- Expectation Maximization
- Contrastive Methods
- Monte Carlo
- Variational Methods

Many more not listed!!!


Families of Generative Models



Learning Principles


● Maximum Likelihood

$$\max_{\theta} \log p_{\theta}(\mathcal{D})$$




E.g. Autoregressive Models

Families of Generative Models



Learning Principles

● Expectation Maximization



E.g. Gaussian Mixture Models

Families of Generative Models



Learning Principles

● Contrastive Methods

E.g. Noise Contrastive Estimation

$$\mathcal{L} = -\log(s(t)) - \log\left(s(t) + \sum_{i=1}^N s(t')\right)$$

Score on positive samples

Score on negative samples

Families of Generative Models



Learning Principles



Monte Carlo

E.g. MCMC

Think Monte-Carlo Gradient Estimation

Families of Generative Models



Learning Principles



Variational Methods

E.g. Amortized Variational Inference

$$\log p(\mathcal{D}) \geq ELBO$$

Families of Generative Models



Algorithms

Combining models + Learning Principles

- LVMs + Variational Inference
- Implicit Generative Model + Two Sample Test
- Autoregressive Model + Maximum Likelihood

Latent Variable Models



Conventional Approaches to Latent Variable Modelling


$$x \in \mathbb{R}^n \quad z \in \mathbb{R}^d \quad \mathcal{D} = \{x_i\} \quad i \in \{1, \dots, N\}$$

Conventional Approaches to Latent Variable Modelling

$$x \in \mathbb{R}^n \quad z \in \mathbb{R}^d \quad \mathcal{D} = \{x_i\} \quad i \in \{1, \dots, N\}$$

$$\log p(x) = \log \int p(x | z) p(z) dz = \log \mathbb{E}_{p(z)}[p(x | z)]$$

$$\log p(\mathcal{D}) = \sum_{i=1}^N \log \mathbb{E}_{p(z)}[p(x_i | z)]$$




Methods for Approximate Inference

- Laplace approximations
- Importance Sampling
- **Variational approximations** → This lecture
- Perturbative corrections
- Other methods: MCMC, Langevin, HMC etc...

Variational Inference

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) dz$$



$p(x | z)$

Variational Inference

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) dz$$

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) \frac{q(z)}{q(z)} dz$$

Importance Sampling

Variational Inference

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) dz$$

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) \frac{q(z)}{q(z)} dz$$

Importance Sampling

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)} \left[\log \frac{p(x_i, z_i)}{q_i(z)} \right]$$

Jensen's Inequality

Variational Inference

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) dz$$

$$\log p(\mathcal{D}) = \log \int p(x | z) p(z) \frac{q(z)}{q(z)} dz$$

Importance Sampling

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)} \left[\log \frac{p(x_i, z_i)}{q_i(z)} \right]$$

Jensen's Inequality

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)} [\log p(x_i | z)] + \mathbb{E}_{q_i(z)} \left[\log \frac{p(z)}{q_i(z)} \right]$$

Variational Inference

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)}[\log p(x_i | z)] + \mathbb{E}_{q_i(z)} \left[\log \frac{p(z)}{q_i(z)} \right]$$

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)}[\log p(x_i | z)] - D_{KL}(q_i(z) || p(z))$$

Evidence Lower Bound (ELBO)

$$\log p(\mathcal{D}) \geq -\mathcal{F}(x, z)$$

Free Energy of the System

Variational Inference

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)}[\log p(x_i | z)] + \mathbb{E}_{q_i(z)} \left[\log \frac{p(z)}{q_i(z)} \right]$$

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_i(z)}[\log p(x_i | z)] - D_{KL}(q_i(z) || p(z))$$

Reconstruction Error

Regularizer

Variational Auto-Encoder

Arguably the catalyst to the Deep Generative Modelling Revolution.


Variational Auto-Encoders

Key Innovations:

- Amortized Inference
- Reparametrization Trick

Amortized Inference

- Introduce a parametric family of densities *shared over all the data*




$$\arg \max_{q_i} \mathbb{E}_{q_i^*(z)} [-\mathcal{F}(x_i, z)] \longrightarrow \arg \max_{\phi} \mathbb{E}_{q_{\phi}(z|x)} [-\mathcal{F}_{\phi}(x_i, z)]$$

VAE - Generative Model

- Sample a Reparametrized Latent z

Latent sample

$$z \sim q_{\phi}(z | x) \longrightarrow$$



Reconstructed datum

$$\longrightarrow \tilde{x}$$


$$\log p(\mathcal{D}) \geq \underbrace{\mathbb{E}_{q_{\phi}(z|x)}[\log p_{\psi}(x | z)]}_{\text{Reconstruction Error}} - \underbrace{D_{KL}(q_{\phi}(z | x) || p(z))}_{\text{Regularizer}}$$

VAE - All the pieces

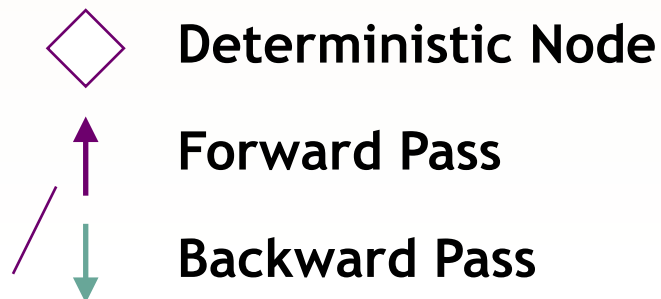
- Prior $p(z)$ usually a standard Normal $\mathcal{N}(0, I)$
- Encoder network $q_\phi(z | x)$
- Sample from the approximate posterior $z \sim q_\phi(z | x)$
- Generative Network $p_\psi(x | z)$
- Model parameters $\theta = \{\phi, \psi\}$
- Objective: Optimize ELBO



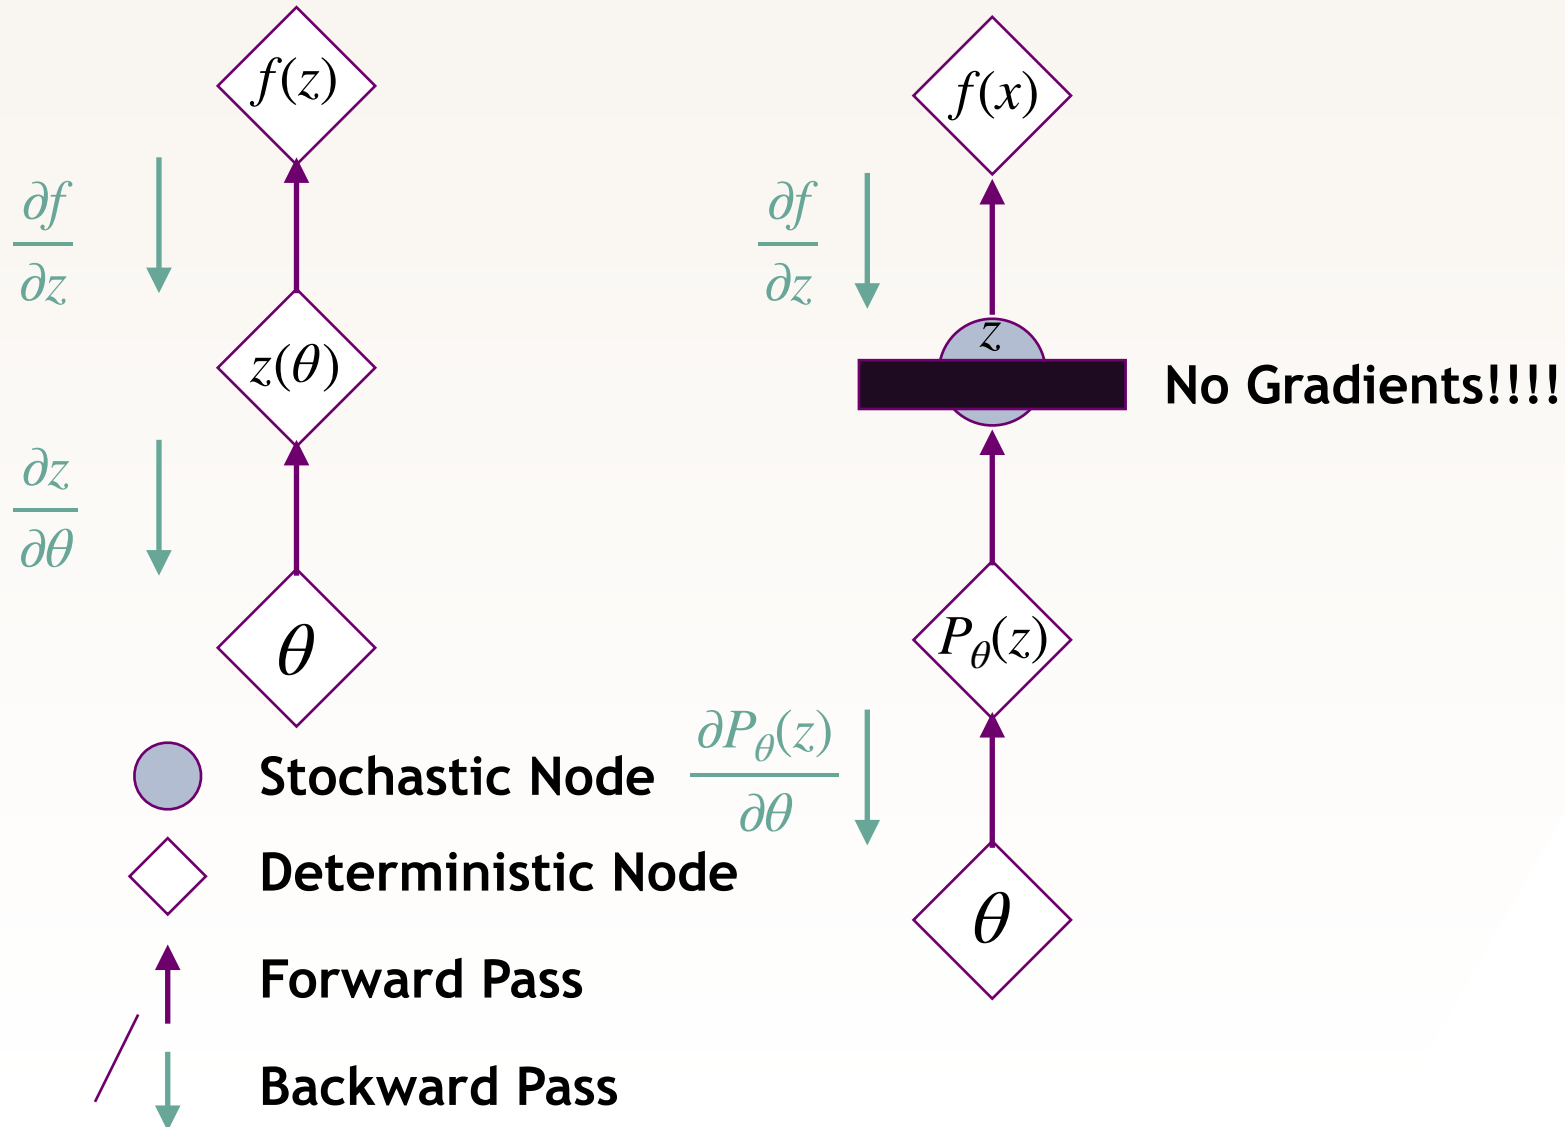
What about Reparametrization?



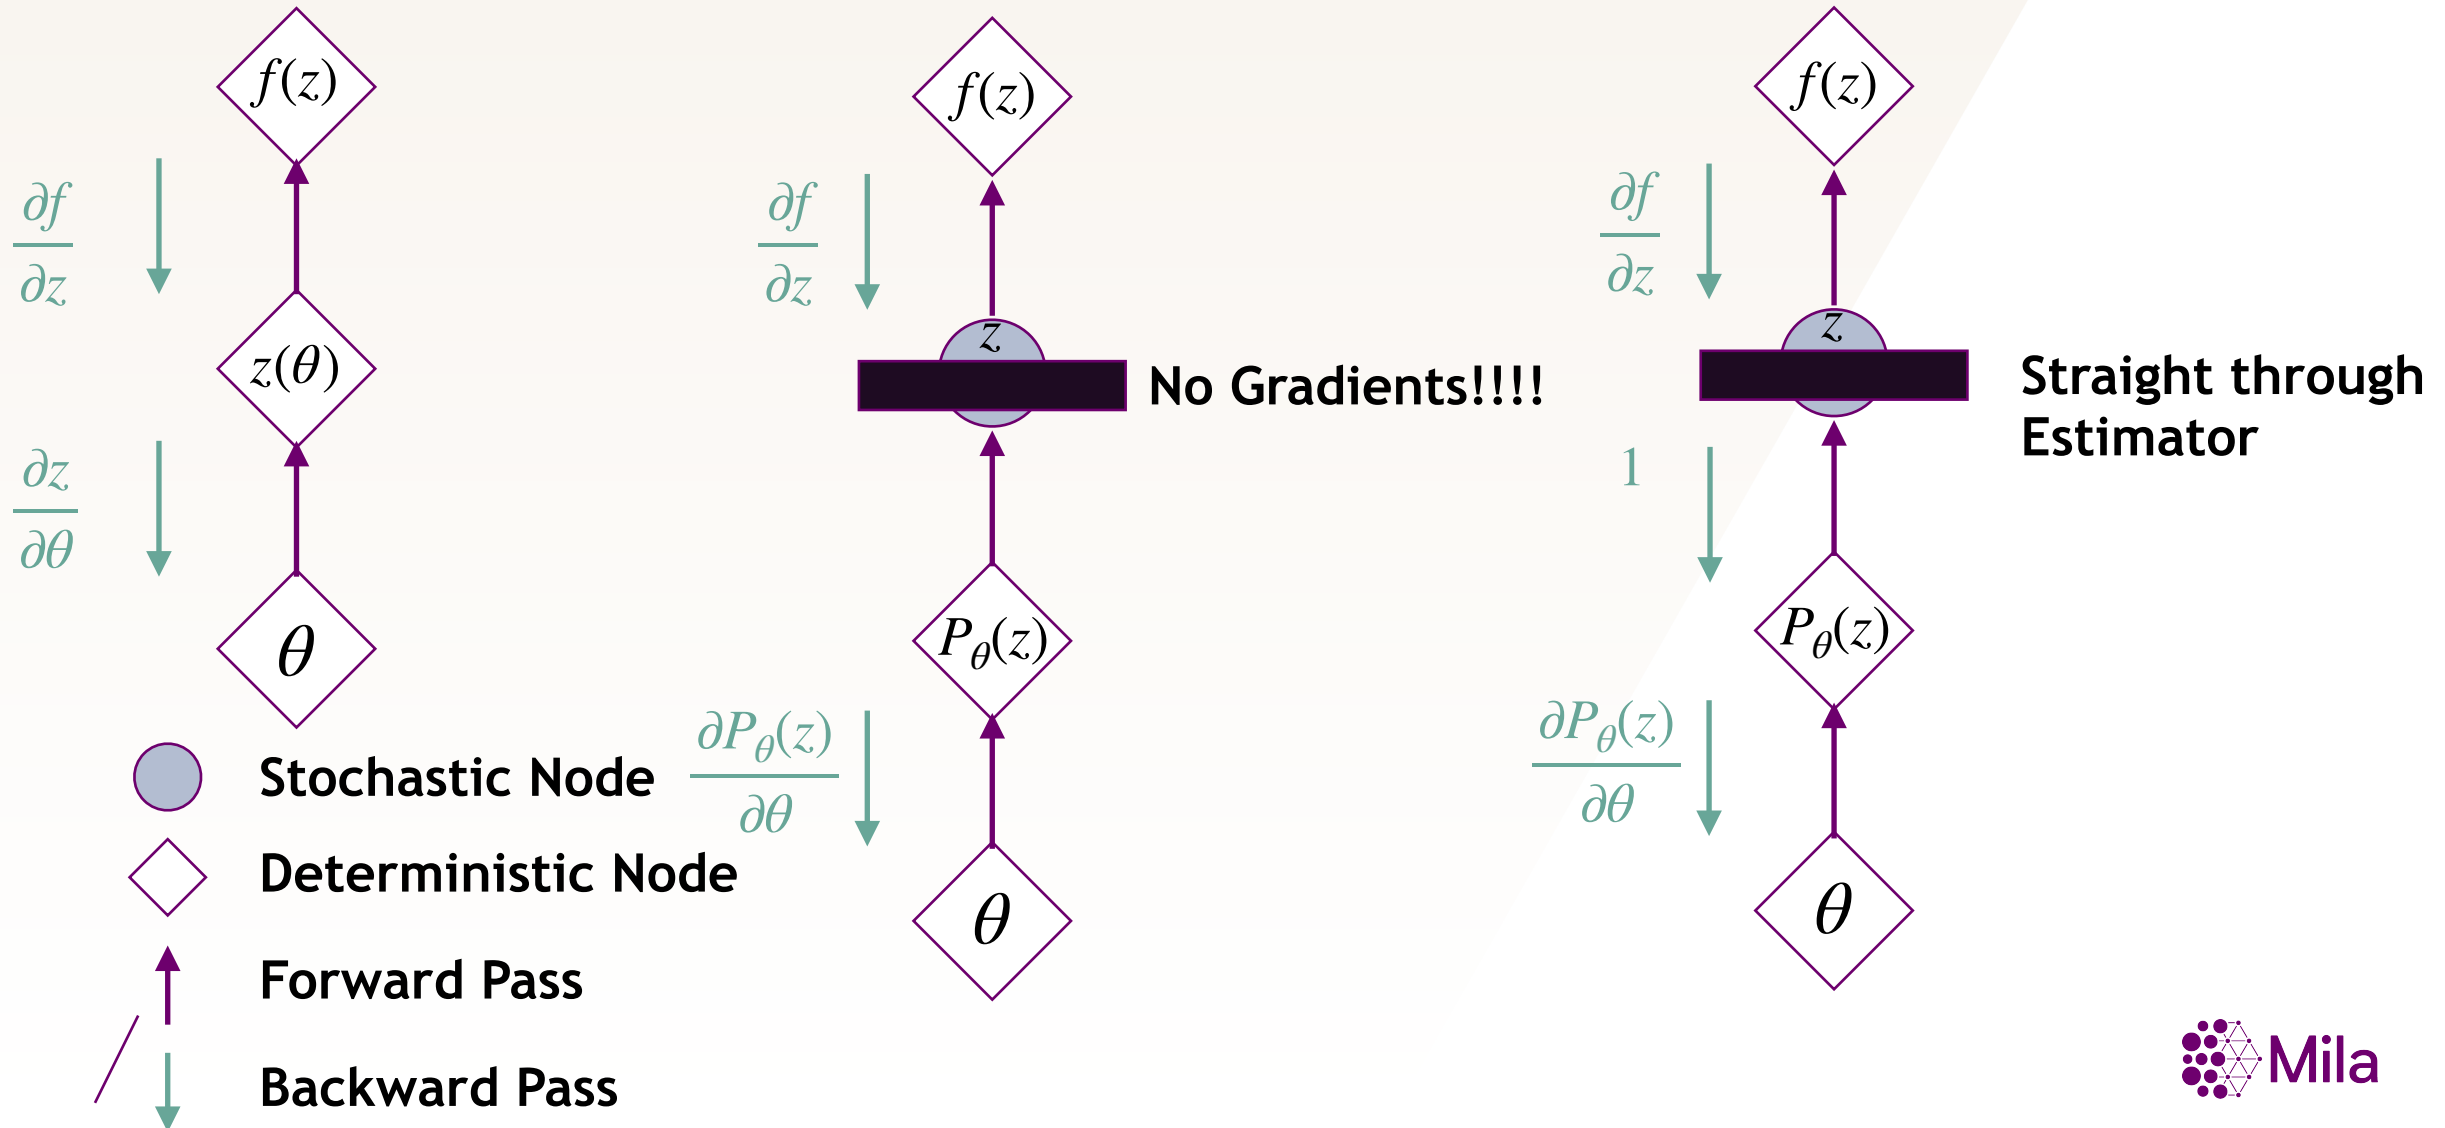
This is the classical setting in many DL problems



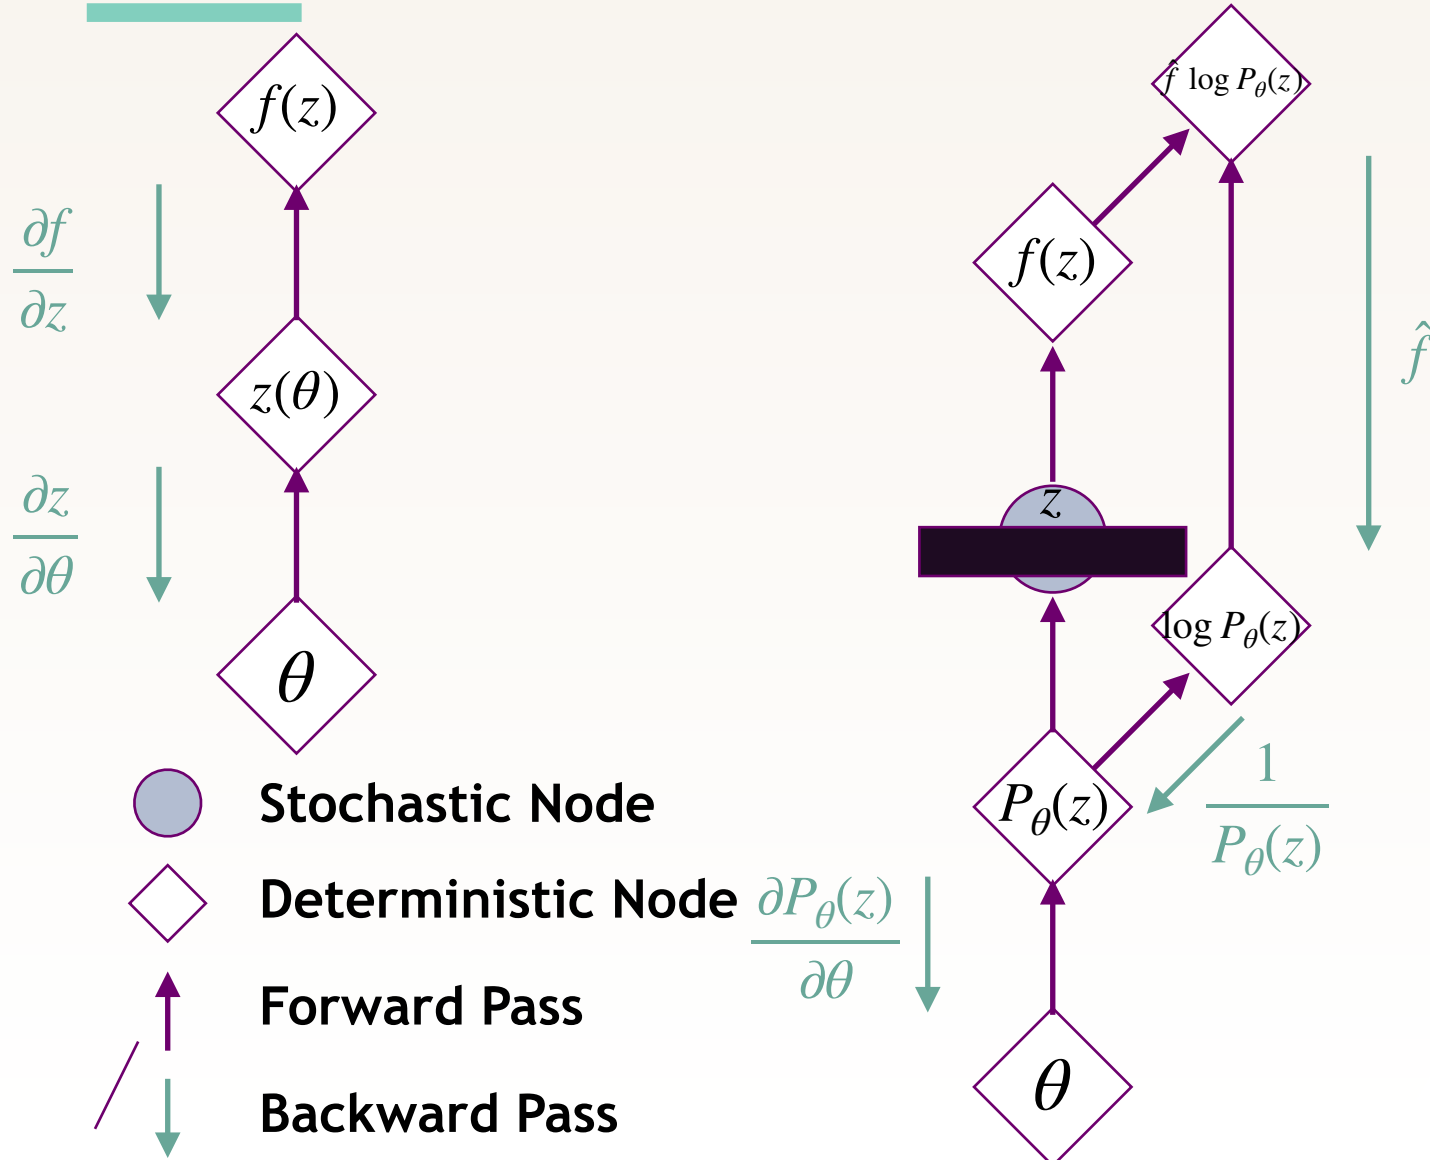
What about Reparametrization?



What about Reparametrization?

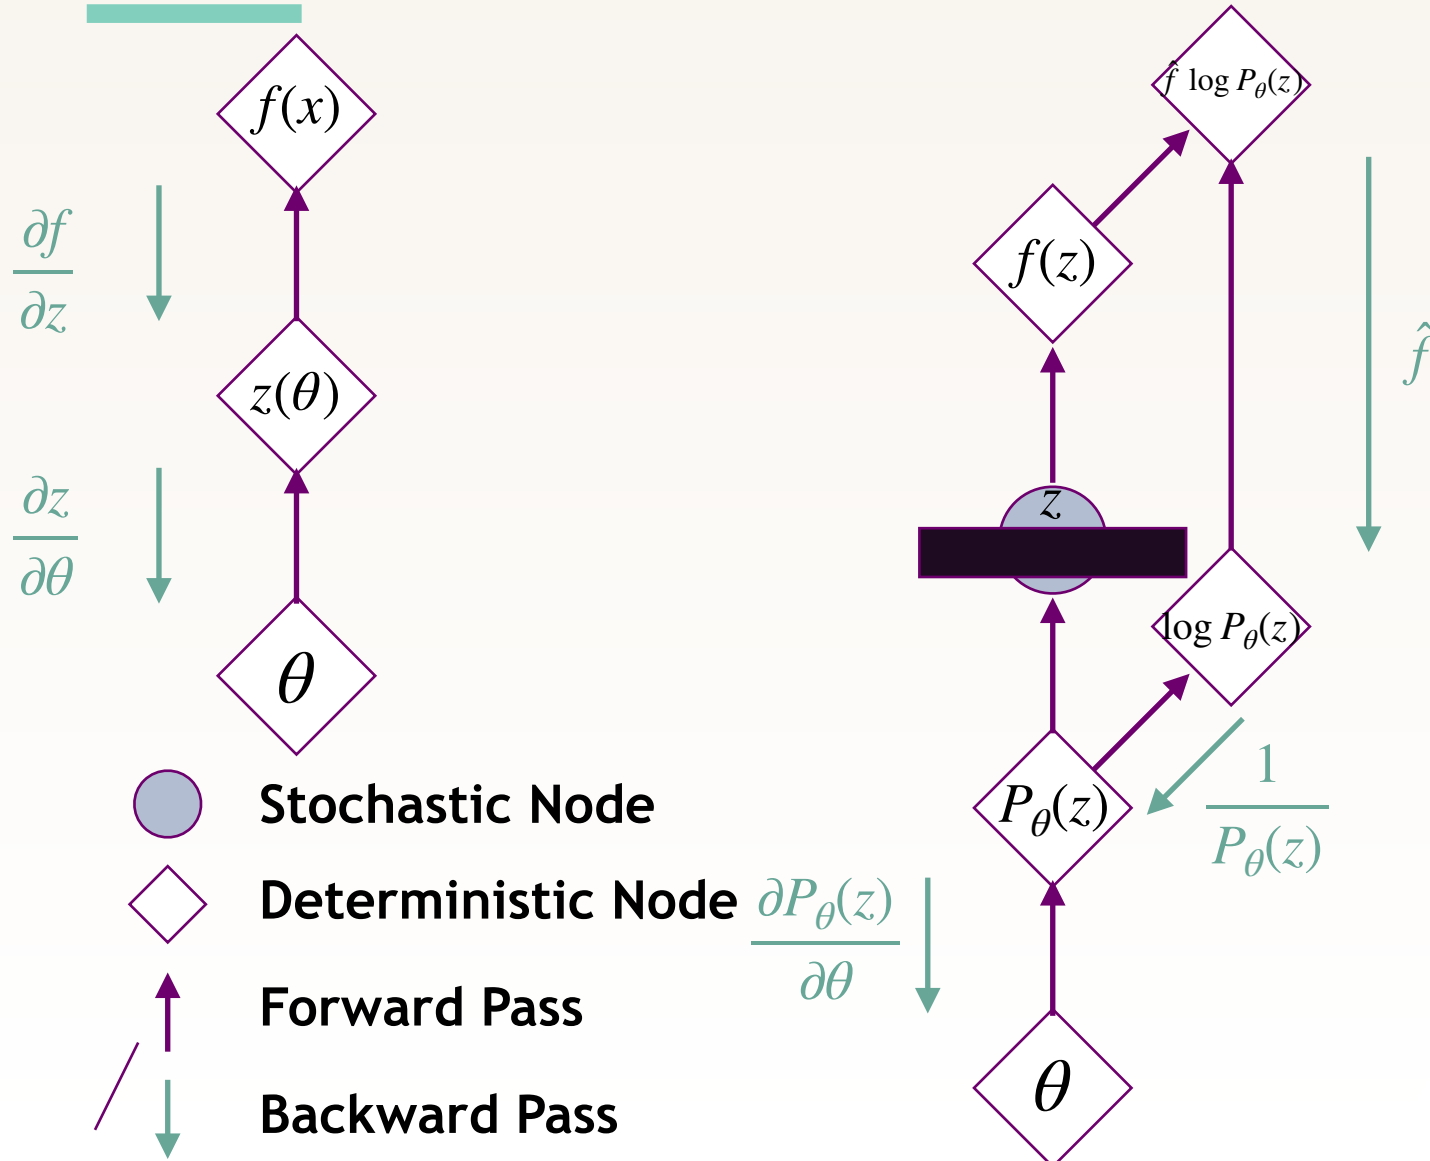


What about Reparametrization?



Does this look familiar to those who do RL?

What about Reparametrization?



Does this look familiar to those who do RL?

Problem:

$$\nabla_{\theta} \mathbb{E}_{z \sim p_{\theta}(z)} [f(z)]$$

REINFORCE:

$$\nabla_{\theta} \mathbb{E}_{z \sim p_{\theta}(z)} [f(z)] = \nabla_{\theta} \mathbb{E}_{z \sim p_{\theta}(z)} [f(z) \nabla_{\theta} \log p_{\theta}(z)]$$

Reparametrization Trick in Gaussian VAE's


$$z \sim q_{\phi}(z|x)$$

$$z = \mu + \epsilon \odot \sigma^2$$

$$z = \mu + \epsilon \odot \sigma^2$$

$$\epsilon \sim \mathcal{N}(0, I)$$

We can backprop normally now!



Application: Lossless Compression

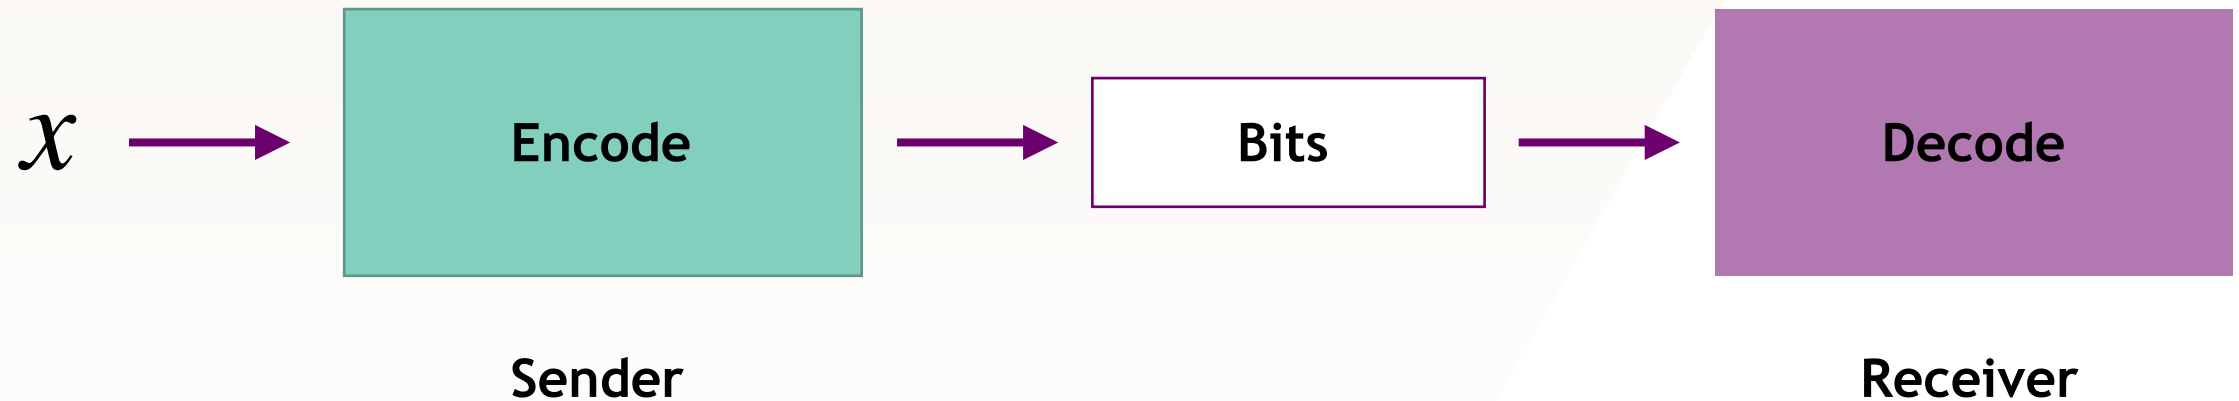
Statistical patterns in the data

“... randomising letters in the middle of words [has] little or no effect on the ability of skilled readers to understand the text. This is easy to demonstrate. In a publication of New Scientist you could randomise all the letters, keeping the first two and last two the same, and readability would hardly be affected. My analysis did not come to much because the theory at the time was for shape and sentence recognition. Saberi's work suggests we may have some powerful parallel processors at work. The reason for this is surely that identifying content by parallel processing speeds up recognition. We only need the first and last two letters to spot changes in meaning.”

--- Graham Rawlinson, PhD Thesis 1976, Nottingham U.

Application: Lossless Compression

Compression is a process that encodes the original data into a representation that takes fewer bits.



Application: Lossless Compression

The Coding Problem:

- Suppose we are given $M : s_1, \dots, s_M$
- We want to create a bit string for these symbols $(0,1)$
- How many bits do we need per symbol?

Without any further information we would need $\log_2 M$ bits



Application: Lossless Compression

The Coding Problem:

- Suppose we are given $M : s_1, \dots, s_M$
- The sender and receiver both know the alphabet \mathcal{A}
- The sender and receiver both have access to a probabilistic model for each symbol.
- Assume Sender and Receiver have agreed upon a standard encoding/decoding paradigm: Arithmetic coding/ANS

How can we communicate the data with the fewest number of bits?



Application: Lossless Compression

Shannon's Source Coding Theorem

$$H(x) = - \mathbb{E}_{p(x)}[\log p(x)]$$

This is the best that we can do we both sender and receiver both use $p(x)$

**In practice the receiver may not have access to $p(z | x)$.
What should they do then? Approximate it!**

Application: Lossless Compression

Entropy Coding Example

$$\mathbf{A} \quad p_A = 0.5$$

$$\mathbf{B} \quad p_B = 0.5$$

$$H(A, B) = -p_A \log_2 p_A - p_B \log_2 p_B$$

$$H(A, B) = -0.5 \log_2 0.5 - 0.5 \log_2 0.5$$

$$H(A, B) = 1$$

$$\mathbf{A} \quad p_A = 0.8$$

$$\mathbf{B} \quad p_B = 0.2$$

$$H(A, B) = -p_A \log_2 p_A - p_B \log_2 p_B$$

$$H(A, B) = -0.8 \log_2 0.8 - 0.2 \log_2 0.2$$

$$H(A, B) \approx 0.7219$$

Less than 1 bit/symbol

Application: Lossless Compression

Some common coding strategies

- Huffman Coding
- Arithmetic Coding
- Lempel-Ziv
- Entropy Coding

**Handling
continuous and
high dimensional
data is a problem!**



Application: Lossless Compression

Notation:

- x - data we want to compress
- z - latent variable; compressed representation
- $p(x)$ - marginal or data distribution
- $p(z | x)$ - the true posterior distribution over data
- $q(z | x)$ - Approximate posterior distribution over data



Application: Lossless Compression

How can the sender communicate $\tilde{x} \sim p(x)$?

- Sender and Receiver both have access to $p(z)$ and $p(x|z)$
- 1. Sender samples $z \sim p(z)$ and then uses it to encode \tilde{x} using $p(\tilde{x}|z)$
- 2. Effective code length is $-(\log p(\tilde{x}) + \log p(\tilde{x}|z))$ bits.

Can we do better?



Application: Bits-Back Coding

Suppose we have access to an LVM

$p(x | z)p(z)$ We compress z first using $p(z)$. Then we compress x using $p(x | z)$.

The receiver would execute this process in reverse to decode

$$H(x, z) = - \mathbb{E}_{p(x, z)} [\log p(x, z)]$$

$$= H(p(x | z)) + H(p(z))$$

Size of bitstream
to encode x

Size of bitstream
to encode z

Application: Bits-Back Coding

Clearly,

$$H(p(x, z)) \geq H(p(x))$$

What's the optimality gap?

Recall $p(x | z) = \frac{p(z | x)p(z)}{p(x)}$

$$\begin{aligned} H(x) &= - \mathbb{E}_{p(x)} \left[\log \frac{p(x | z)p(z)}{p(z | x)} \right] \\ &= - \mathbb{E}_{p(x, z)} [\log p(x | z)] - \mathbb{E}_{p(x, z)} [\log p(z)] + \mathbb{E}_{p(x, z)} [\log p(z | x)] \\ &= H(p(x | z)) + H(p(z)) - H(p(z | x)) \end{aligned}$$

/ We overestimate the length by $H(p(z | x))$ bits, Bits-Back will address this

Application: Bits-Back Coding

But the receiver doesn't have $p(z | x)$?

$q(z | x)$ We need to build an approximate posterior.

What's the minimum code length now?

$$\begin{aligned} & \mathbb{E}_{q(z|x)} \left[-\log p(x | z) - \log p(z) + \log q(z | x) \right] \\ &= \mathbb{E}_{q(z|x)} \left[-\log p(x, z) + \log q(z | x) \right] \\ &= \mathbb{E}_{q(z|x)} \left[-\log p(z | x) - \log p(x) + \log q(z | x) \right] \\ &= \mathbb{E}_{q(z|x)} \left[-\log p(x) + \log \frac{q(z | x)}{p(z | x)} \right] \end{aligned}$$

Application: Bits-Back Coding

But the receiver doesn't have $p(z | x)$?

$q(z | x)$ We need to build an approximate posterior.

What's the minimum code length now?

$$= \mathbb{E}_{q(z|x)} \left[-\log p(x) + \log \frac{q(z|x)}{p(z|x)} \right]$$

$$= \mathbb{E}_{q(z|x)} [-\log p(x)] + D_{KL}(q(z|x) || p(z|x))$$

$$\log p(x) = ELBO + D_{KL}(q(z|x) || p(z|x)) \quad \text{KL-Gap due to mismatch between posteriors}$$

$$\geq 0$$

Application: Bits-Back Coding

Suppose the Sender has additional bits (possibly random)

1. Sender samples uses random bits to generate a sample \tilde{z} by using $p(z|x)$ to decode! Generating this sample costs $-\log p(\tilde{z})$ bits.
2. Sender uses \tilde{z} to encode \tilde{x} using $p(\tilde{x}|\tilde{z})$ and also encodes \tilde{z} using $p(\tilde{z})$. This costs $-(\log p(\tilde{x}) + \log p(\tilde{x}|\tilde{z}))$ as before.
3. The Receiver can now recover this junk information by first decoding \tilde{x} then \tilde{z} , effectively getting the “bits-back”. The net cost for this scheme is $-\log p(\tilde{x}|\tilde{z}) - \log p(\tilde{z}) + \log p(\tilde{z}) = -\log p(\tilde{x}|\tilde{z})$.



Application: Bits-Back Coding

High Level Idea:

- Sender will use $p(x | z)$, $p(z)$, and $p(z | x)$ as opposed to the first two in LVM.
- The goal is to annihilate the overestimation gap $-H(p(z | x))$.
- If Encoder sees terms from $H(p(x | z))$ or $H(p(z))$ we push to the bitstream making it longer. If decoder sees these terms we pop.
- If Encoder sees terms from $-H(p(z | x))$ we pop. Conversely, if decoder sees these terms we push to the bitstream.
- During encoding we pop $-H(p(z | x))$ and push $H(p(x | z))$ and $H(p(z))$. Same for Decoding.

Application: Bits-Back Coding

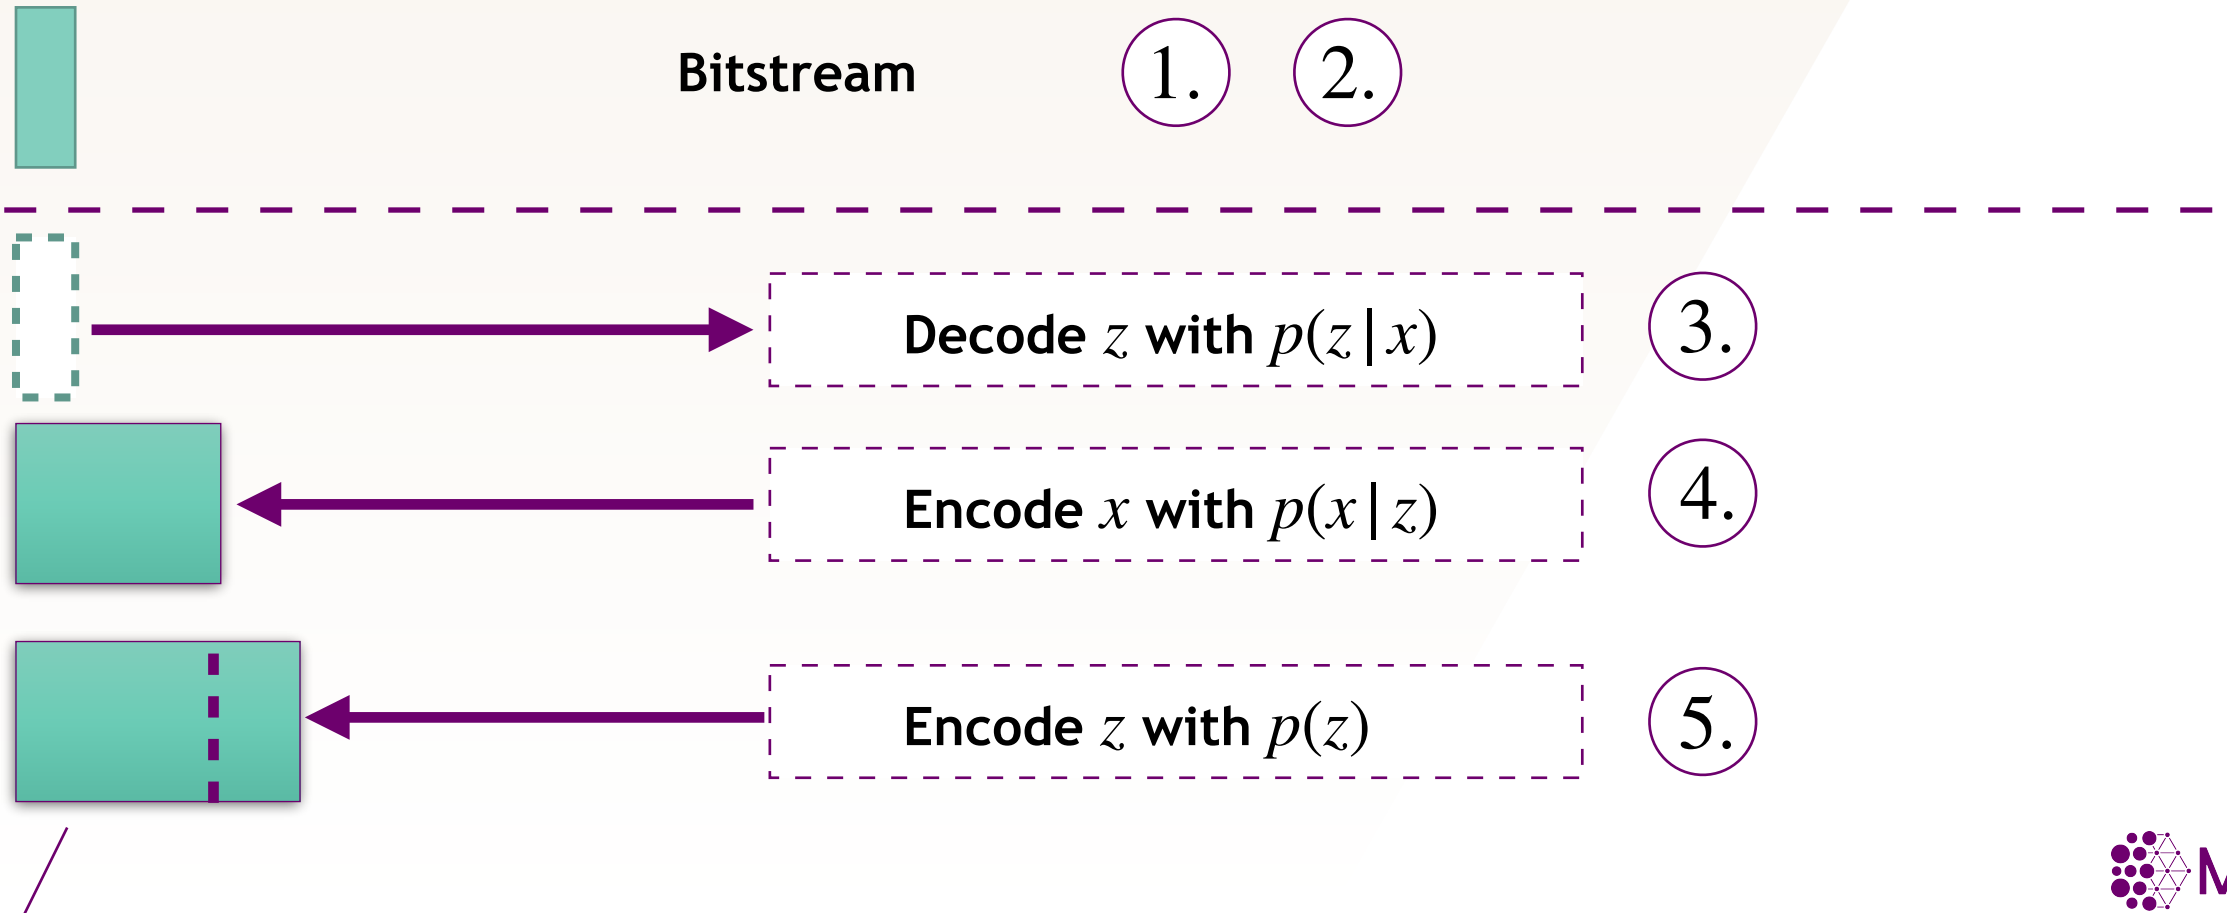
Sender Perspective

1. Assume there is an initial Bitstream
2. Sender sends x ; has access to $x, z, p(x|z), p(z), p(z|x)$
3. Sender uses initial bit-stream and decodes z using $p(z|x)$. Pop from bitstream
4. Sender encodes x using $p(x|z)$. Push onto bitstream.
5. Sender encodes z using $p(z)$. Push onto bitstream.



Application: Bits-Back Coding

Sender Perspective



Application: Bits-Back Coding

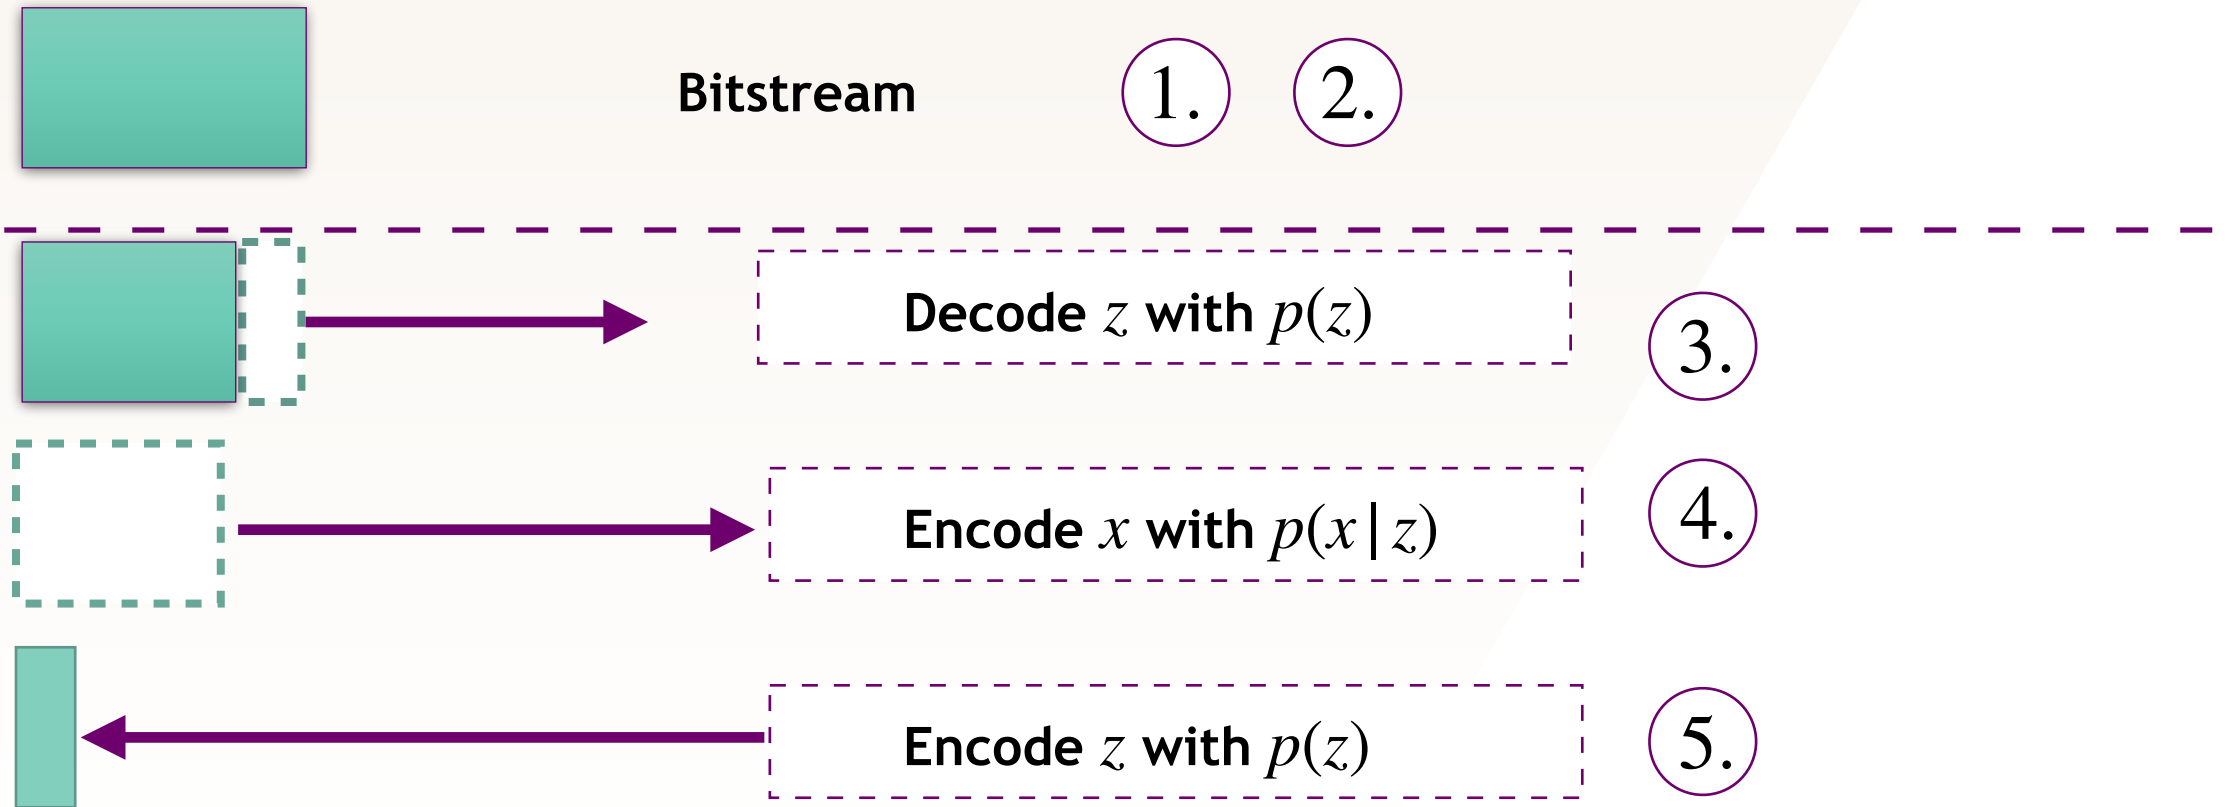
Receiver Perspective

1. Receiver has access to $p(x | z), p(z), p(z | x)$.
2. Receiver decodes z using $p(z)$. Pop from bitstream.
3. Receiver decodes x using $p(x | z)$. Pop from bitstream.
4. Receiver encodes z using $p(z | x)$. Push to the bitstream.
5. Receiver successfully decodes x and recovers the initial bitstream.



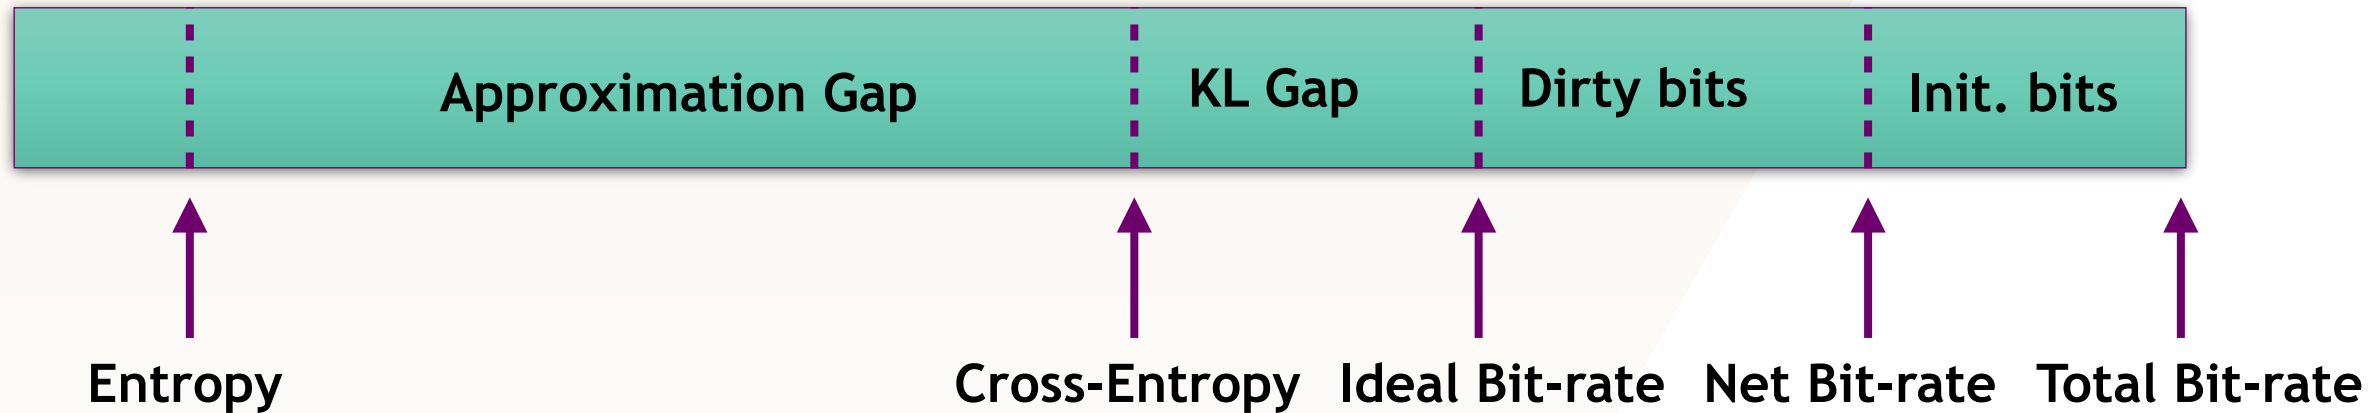
Application: Bits-Back Coding

Receiver Perspective



Application: Bits-Back Coding

The total Bitstream



Addition VAE Topics Not Covered

- Latent Space Disentanglement
- Hierarchical Latent Spaces
- Non-Gaussian Priors
- Identifiability and Causality



GANs

Game Theorists are still thanking Ian Goodfellow et. al 2014 for making their field profitable.

Implicit Generative Models

- What if we design a model where we cannot exactly compute the density?

We can no longer train using Maximum Likelihood and ELBO!



Implicit Generative Models

- What if we design a model where we cannot exactly compute the density?

We can no longer train using Maximum Likelihood and ELBO!

- We can use a comparison metric between our model density q and p^*

This is known as Implicit Density Estimation



Implicit Generative Models

- We can use the Density Ratio as our Comparison Metric: $r = \frac{p^*(x)}{q(x)}$

Why is this sensible?

- Suppose we had class labels $y \in [-1, 1]$, indicating Fake vs. Real samples

$$r = \frac{p^*(x)}{q(x)}$$

Implicit Generative Models

- Suppose we had class labels $y \in [-1, 1]$, indicating Fake vs. Real samples

$$r = \frac{p^*(x)}{q(x)}$$

$$r = \frac{p^*(x | y = 1)}{q^*(x | y = -1)}$$

Implicit Generative Models

- Suppose we had class labels $y \in [-1, 1]$, indicating Fake vs. Real samples

$$r = \frac{p^*(x)}{q(x)}$$

$$r = \frac{p^*(x | y = 1)}{q(x | y = -1)}$$

$$r = \frac{p^*(y = 1 | x)}{q(y = -1 | x)}$$

Bayes Rule + Cancellations

Implicit Generative Models

- Suppose we had class labels $y \in [-1, 1]$, indicating Fake vs. Real samples

$$r = \frac{p^*(y = 1 | x)}{q(y = -1 | x)} \quad \text{Bayes Rule + Cancellations}$$

- This is more manageable, and we can introduce a parametric scoring function


$$p^*(y = 1 | x) = D_{\theta}(x)$$

$$p^*(y = -1 | x) = 1 - D_{\theta}(x)$$

Generative Adversarial Networks

Latent sample

$$z \sim p(z)$$

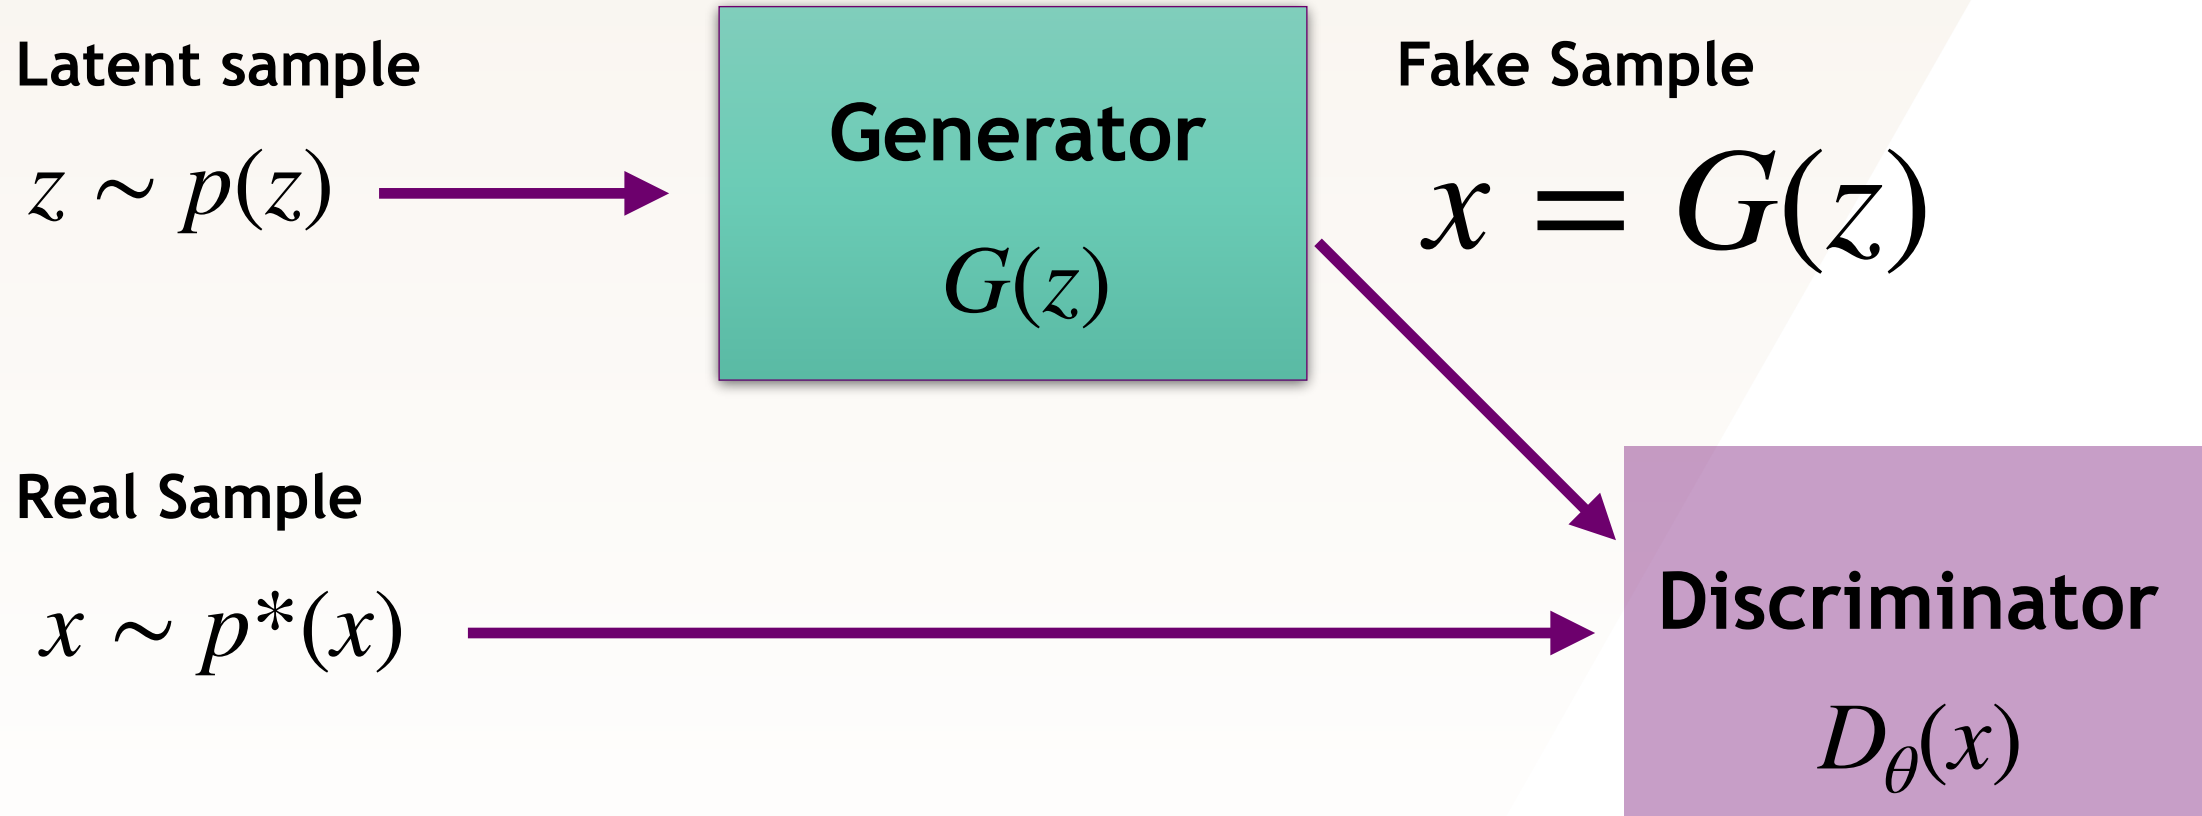


Fake Sample




$$x = G(z)$$

Generative Adversarial Networks



Generative Adversarial Networks



$$\min_G \max_D \mathbb{E}_{p^*(x)}[\log(D_{\theta}(x))] + \mathbb{E}_{q_{\phi}(\tilde{x})}[\log(1 - D_{\theta}(\tilde{x}))]$$

Generative Adversarial Networks

Discriminator Loss

$$\mathcal{L}_D = \mathbb{E}_{p^*(x)}[-\log(D_\theta(x))] + \mathbb{E}_z[\log(1 - D_\theta(G_\phi(z)))]$$

Generator Loss

$$\mathcal{L}_G = -\mathcal{L}_D$$

Combined Loss

$$\min_G \max_D \mathbb{E}_{p^*(x)}[\log(D_\theta(x))] + \mathbb{E}_{q_\phi(\tilde{x})}[\log(1 - D_\theta(\tilde{x}))]$$

Generative Adversarial Networks

Discriminator Loss

$$\mathcal{L}_D = \mathbb{E}_{p^*(x)}[-\log(D_\theta(x))] + \mathbb{E}_z[\log(1 - D_\theta(G_\phi(z)))]$$

Generator Loss

$$\mathcal{L}_G = -\mathcal{L}_D$$

Non-Saturating Generator Loss

$$\mathcal{L}_G = -\mathbb{E}_z[\log D_\theta(G(z))]$$

Combined Loss

$$\min_G \max_D \mathbb{E}_{p^*(x)}[\log(D_\theta(x))] + \mathbb{E}_{q_\phi(\tilde{x})}[\log(1 - D_\theta(\tilde{x}))]$$

Many Breeds of GANs

- Wasserstein GAN
- Wasserstein GAN with Gradient Penalty/Spectral Norm
- CycleGAN
- StyleGAN
- BigGAN



Normalizing Flows

Invented at Mila by Laurent Dinh et. al and at Deep Mind by Danilo Rezende et. al in parallel.

Motivation for Normalizing Flows

- What if we want to do exact Maximum Likelihood density estimation?
- What if we also want fast sampling as opposed to autoregressive models?
- What if we also want to have a rich family of models that are easier to train than GANs?

Setup Normalizing Flows

$$x \in \mathbb{R}^n \quad z \in \mathbb{R}^n \quad \mathcal{D} = \{x_i\} \quad i \in \{1, \dots, N\}$$

Take a bijective Function

$$f: \mathbb{R}^n \rightarrow \mathbb{R}^n \quad f^{-1} = g \quad g \circ f(z) = z$$

How does the random variable z with distribution $q(z)$ transform under f ?

$$z' = f(z)$$

Setup Normalizing Flows

How does the random variable z with distribution $q(z)$ transform under f ?

$$z' = f(z)$$

$$q(z') = q(z) \left| \det \frac{\partial f^{-1}}{\partial z'} \right|$$


**Change of variable formula
for Probability Distribution**

$$q(z') = q(z) \left| \det \frac{\partial f}{\partial z} \right|^{-1}$$

Inverse function theorem

1D Example


Let X be Uniform(0,1). Let $Y = f(X) = 2X + 1$



Observe that the probability mass must integrate to 1.

1D Example


Let X be Uniform(0,1). Let $Y = f(X) = 2X + 1$



Observe that the probability mass must integrate to 1.

1D Example

Let X be Uniform(0,1). Let $Y = f(X) = 2X + 1$



Area of a parallelogram is $ad - bc$.

This is nothing but the determinant of a 2×2 matrix.

Density Estimation with Normalizing Flows

- Each function f_i must be invertible.
- We must be able to efficiently sample from the final distribution,
 $z_k = f_k \circ f_{k-1} \circ \dots \circ f_1(z_0)$
- We must be able to efficiently compute the associated change in volume

Change in Volume

$$\log p(z_k) = \log p(z_0) - \sum_{j=1}^k \log \det \left| \frac{\partial f_j}{\partial z_{j-1}} \right|$$

Affine Coupling Flows

Forward Transform

$$y_{1:d} = x_{1:d}$$

$$y_{d+1:D} = x_{d+1:D} \odot \exp(s(x_{1:d})) + t(x_{1:d})$$

Inverse Transform


$$x_{1:d} = y_{1:d}$$

$$x_{d+1:D} = (y_{d+1:D} - t(y_{1:d}) \odot \exp(-s(y_{1:d})))$$


$$\frac{\partial y}{\partial x} = \begin{bmatrix} \mathbb{I}_d & 0 \\ \frac{\partial y_{d+1:D}}{\partial x_{1:d}^T} & \text{diag}(\exp(s(x_{1:d}))) \end{bmatrix}.$$

Affine Coupling Flows


Forward Transform




Inverse Transform



Residual Flows



Residual Flows



Forward Dynamics

$$x_{t+1} \longleftarrow x_t + \alpha h(x_t)$$

This looks an Euler Discretization of an ODE.

Reverse Dynamics

$$x_t \longleftarrow x_{t+1} - \alpha h(x_t)$$

When can we invert a ResNet?

Residual Flows

Sufficient Condition

$$\text{Lip}(h_t) \leq 1$$

This condition basically tells us the IVP has a solution that not only exists and is unique. See Picard-Lindelof Theorem.

Forward Dynamics

$$x_{t+1} \longleftarrow x_t + \alpha h(x_t)$$

This looks an Euler Discretization of an ODE

Reverse Dynamics

$$x_t \longleftarrow x_{t+1} - \alpha h(x_t)$$

When can we invert a ResNet?

Residual Flows

Change of Variable

$$f(x) = h(x) + x$$

$$\log p(x) = \log p(f(x)) + \log \left| \det \frac{df(x)}{dx} \right|$$

$$\log p(x) = \log p(f(x)) + \text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_h(x)]^k \right) \quad J_h = \frac{dh(x)}{dx}$$

Where did this come from?



Residual Flows

Change of Variable

$$\text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_h(x)]^k \right) \quad J_h \equiv \frac{df(x)}{dx}$$

$$\left| \det J_f \right| = \det J_f$$

Lipschitz constrained perturbations of the form $x + h(x)$ have positive det



Residual Flows

Change of Variable

$$\text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_h(x)]^k \right) \quad J_h = \frac{dh(x)}{dx}$$

$$\left| \det J_f \right| = \det J_f$$

Lipschitz constrained perturbations of the form $x + h(x)$ have positive det

$$\ln \det(J_f) = \text{tr}(\ln(J_f)) \quad \ln \det(A) = \text{tr}(\ln(A)) \quad \text{For non-singular matrices } A$$



Matrix Trace and Matrix Logarithm

Residual Flows

Change of Variable

$$\text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_h(x)]^k \right) \quad J_h = \frac{dh(x)}{dx}$$

$$|\det J_f| = \det J_f$$

Lipschitz constrained perturbations of the form $x + h(x)$ have positive det

$$\ln \det(J_f) = \text{tr}(\ln(J_f)) \quad \ln \det(A) = \text{tr}(\ln(A)) \quad \text{For non-singular matrices } A$$

$$\text{tr}(\ln(I + J_h(x))) = \text{tr} \left(\sum_{k=1}^{\infty} \frac{(-1)^{k+1}}{k} [J_h(x)]^k \right) \quad \text{The matrix logarithm has a convergent power series for } \|J_h\|_2 \leq 1$$

Residual Flows

How do we compute the Trace efficiently?

$\text{tr}(J_h)$ Naive computation takes $O(n^2)$

Trick #1

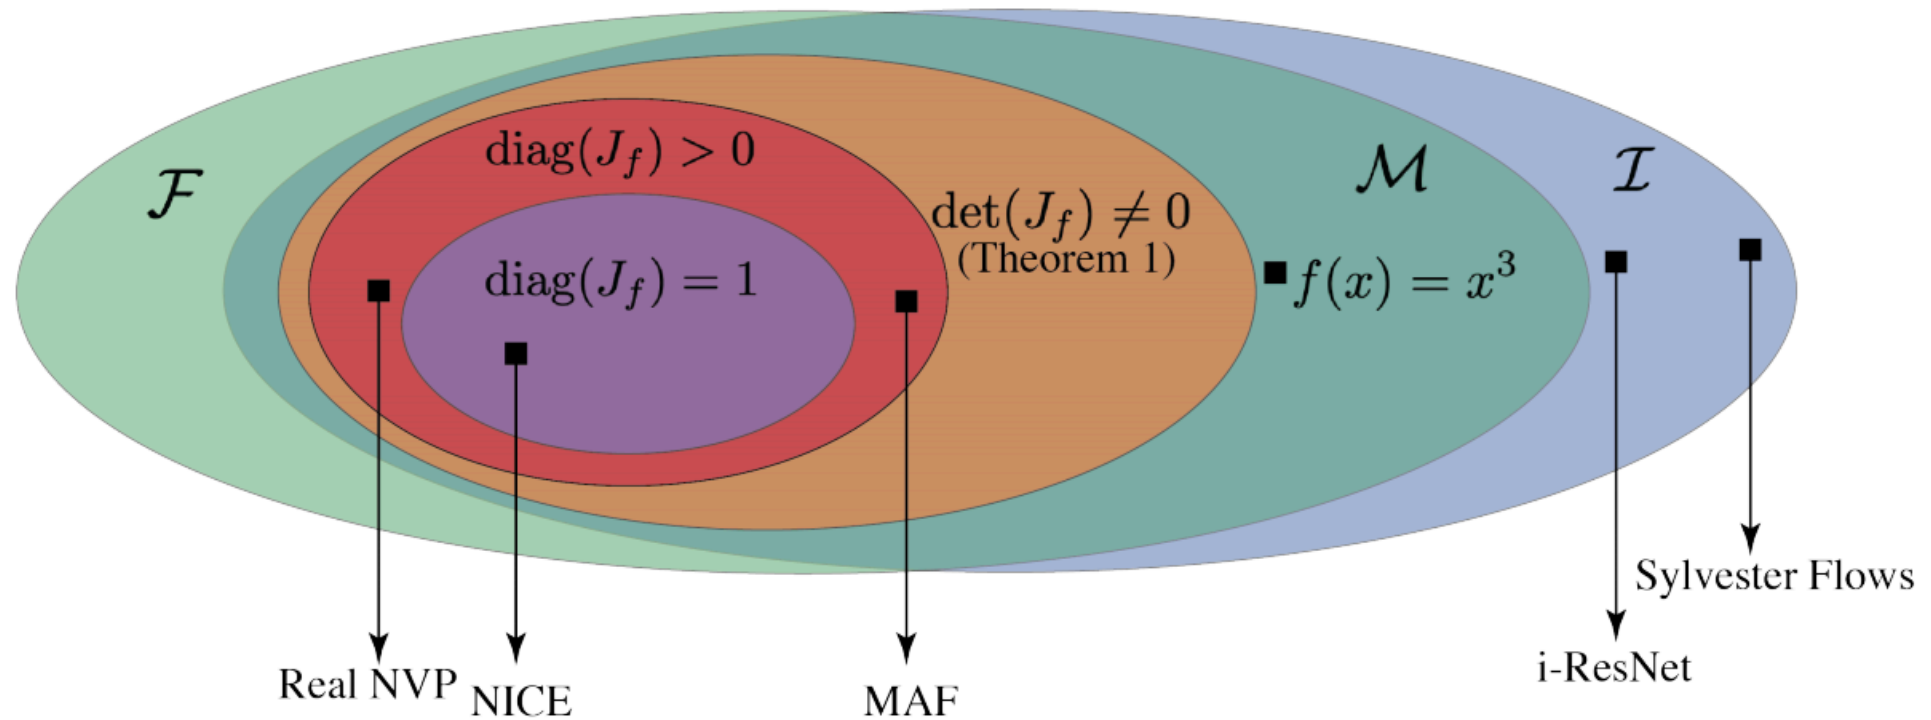
Vector-Jacobian products $v^T J_h$ can be computed at the same cost as a forward pass through h using reverse mode automatic differentiation.

Trick #2

We can compute a stochastic approximation of the trace using the Hutchinson's Trace Estimator

$$\text{tr}(J_h) = \mathbb{E}_{v \sim p(v)}[v^T J_h v] \qquad \mathbb{E}[v] = 0 \qquad \text{Cov}(v) = I$$

Different Normalizing Flows at a glance



Continuous Normalizing Flows

This is the framework that is most amenable to Manifold data.

Neural ODE

Forward Dynamics

$$x_{t+1} \longleftarrow x_t + \alpha h(x_t)$$

This looks an Euler Discretization of an ODE

Reverse Dynamics

$$x_t \longleftarrow x_{t+1} - \alpha h(x_t)$$

When can we invert a ResNet?

What if we take the discretization to be infinitesimal?

$$z_{t_0} = z$$

$$z_{t_1} = z$$

$$\frac{dz_t}{dt} = h_t(t, z_t)$$

This is an ODE!

Continuous Normalizing Flows

Forward Dynamics

$$x = z_{t_1} = z_0 + \int_{t=t_0}^{t_1} h(z_t, t) dt$$

What if we take the discretization to be infinitesimal?

$$z_{t_0} = z$$

$$z_{t_1} = z$$

$$\frac{dz_t}{dt} = h_t(t, z_t)$$

Reverse Dynamics

$$z_0 = x + \int_{t=t_1}^{t_0} h_\phi(z_t, t) dt = x - \int_{t=t_0}^{t_1} h_\phi(z_t, t) dt$$

This is an ODE!

Residual Flows

Forward Dynamics

$$x = z_{t_1} = z_0 + \int_{t=t_0}^{t_1} h(z_t, t) dt$$


Instantaneous Change of Variable

$$\frac{d \log p(z_t)}{dt} = -\text{Tr} \left(J_{h(t, \cdot)}(z_t) \right)$$


Reverse Dynamics

$$z_0 = x + \int_{t=t_1}^{t_0} h_\phi(z_t, t) dt = x - \int_{t=t_0}^{t_1} h_\phi(z_t, t) dt$$

Residual Network



ODE Network



Continuous Normalizing Flows

Instantaneous Change of Variable

$$\frac{d \log p(z_t)}{dt} = - \text{Tr} \left(J_{h(t, \cdot)}(z_t) \right)$$

Change of Variable

$$\log p(x) = \log p(z_0) - \int_{t=t_0}^{t_1} \text{Tr} \left(J_{h_\phi(t, \cdot)}(z_t) \right)$$

Continuous Normalizing Flows

Numerically Solving ODE

- Discrete Solvers that use Euler method e.g. Runge-Kutta
- Adjoint Method

Con's to using CNF's

- CNF's must be globally Lipschitz; this limits their expressive power
- ODE Solvers are slower as solutions must be to a pre-specified tolerance
- Prone to Numerical Errors and Noisy Gradients

Variational Inference with Normalizing Flows

Game Theorists are still thanking Ian Goodfellow et. al 2014 for making their field profitable.

ELBO in VAE

$$\log p(\mathcal{D}) \geq \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\psi}(x|z)] - D_{KL}(q_{\phi}(z|x) || p(z))$$

Reconstruction Error

Regularizer

Sources for a loose ELBO?

- Amortization cost
- Gaussian Posterior Approximation

Improving over the Gaussian Posterior Approximation

We can try to replace the Gaussian Posterior with a Flow

$$\begin{aligned}\mathcal{F}(x) &= \mathbb{E}_{q_0(z_0)}[\log q_k(z_j) - \log p(x, z_j)] \\ &= \mathbb{E}_{q_0(z_0)} \left[\log q_0(z_0) - \sum_{i=1}^j \ln \det \left| \frac{\partial f_i}{\partial z_{i-1}} \right| - \log p(x, z_i) \right] \\ &= D_{KL}(q_0(z_0) || p(z_j)) - \mathbb{E}_{q_0(z_0)} \left[\sum_{i=1}^j \ln \det \left| \frac{\partial f_i}{\partial z_{i-1}} \right| - \log p(x | z_i) \right]\end{aligned}$$